

Rising to the demands of the AI era

Enterprise AI has moved past experimentation. The organizations deploying it at production scale are seeing returns that would have seemed implausible two years ago and that keep accelerating as model capabilities evolve.

A global insurer compressed its underwriting review timeline by more than 5x while improving data accuracy by 75 to 90%. A pharmaceutical company cut clinical study report production from ten weeks to under ten minutes. A telecom provider deployed more than 13,000 custom AI solutions across its workforce in under a year.


These are production deployments running on real data, at scale. And as frontier models and technologies become increasingly more capable, the door to even more advanced and impactful use cases is wide open.

Prior enterprise technology cycles took years to mature, but AI is compressing that timeline into months or even weeks. The organizations that successfully build and deploy these capabilities now accumulate compounding advantages that become increasingly difficult for late movers to close.

Many enterprise AI programs haven't reached this point. According to Accenture's July 2026 **Pulse of Change report**, only 23% of C-suite leaders report having achieved sustained, enterprise-wide AI impact with AI initiatives. Readiness to scale shows a similar shortfall: 64% of respondents to Accenture's May 2026 **AI-Ready Data for Advanced AI survey** say they have moved beyond pilots into production across multiple functions or initiated enterprise-wide efforts for advanced AI, yet only 7% have reached the data-readiness required to scale it. Moving successfully from pilot to an organization-wide rollout is where many programs stall.

Closing the gap between pilot and production

Pilots often succeed because they aren't bound by the same rigorous criteria of production-scale deployments. They operate with protected budgets, purpose-built teams, curated datasets, and highly technical or AI-native engineers who are not representative of the broader enterprise workforce.



Percentage of C-suite leaders who have achieved sustained, enterprise-wide AI impact

Accenture, Pulse of Change, July 2026

Ambition has outrun readiness



Have moved beyond pilots into production across multiple functions, or initiated enterprise-wide efforts for advanced AI



Have reached the data-readiness required to scale it

Accenture, AI-Ready Data for Advanced AI survey, May 2026

They also have limited exposure to the stakeholders and departments that will eventually own whatever gets deployed.

That insulation makes pilots successful, but it also makes them an unreliable signal for how enterprise AI will perform in production at scale. The fix is to incorporate the governance, operational ownership, delivery stability, and economics as a part of the pilot design itself.

That means:

- Defining production success criteria upfront
- Testing against representative enterprise conditions
- Minimizing hidden manual intervention
- Validating delivery scalability beyond AI-native specialist teams
- Embedding governance, observability, and operational controls early
- Measuring long-term operational economics, not just short-term model performance

As a result, organizations run a pilot that functions as a production readiness exercise rather than just a demonstration of technology capabilities.

This guide outlines seven considerations that can help AI programs successfully reach production. Each consideration highlights the operational, organizational, and technical decisions that must be addressed before moving forward.

Each consideration also includes two types of questions:

- **Work out** questions require cross-functional input and have multiple owners
- **Assign** questions are ownership decisions that a CIO or business leader must make

The organizations that get AI to meaningful production treat their pilots with the same diligence they apply to real-world deployments, from success criteria to stakeholder involvement and infrastructure decisions. The seven considerations that follow are drawn from what we've observed across the organizations that made it to production: Anthropic through enterprise deployments built on Claude, Accenture through the implementations it has guided across industries and geographies.

Consideration 01

PRE-PILOT

Strategic
objectives, ROI,
and ownership

Build your evaluation framework before the pressure builds

AI pilots tend to leave success criteria loosely defined, for understandable reasons. The goal of a pilot is to demonstrate that a technology can do something useful, not to specify the exact conditions under which it should be deployed at scale.

The problem is that many pilots operate under artificially optimized conditions: curated data, AI-native engineering teams, narrow scopes, protected budgets, and a high level of manual intervention. Those conditions rarely represent the realities of enterprise production environments.

As organizations move from experimentation to deployment, ambiguity around success criteria becomes a major risk. The questions shift from “Can we model work?” to:

- Can it operate reliably across fragmented enterprise systems and data?
- Can standard delivery teams support it at scale?
- What governance and oversight are required?
- Are the economics sustainable in steady-state operations?

Set the performance bar upfront

When organizations haven't established and agreed on success criteria before production, different stakeholders arrive with different definitions. For example, in an AI-assisted contract review deployment, success might mean something different for each stakeholder:

- The Engineering team wants to track if the system flags non-standard clauses accurately enough that attorney review time drops below what manual review required.
- The Finance team wants to track if the cost per reviewed contract, accounting for attorney time on flagged items, is lower than the fully-manual process.
- The Legal/Risk team wants to track if every output carries a complete audit trail that satisfies their documentation requirements.

Success criteria that don't converge pre-pilot are unlikely to converge post-launch. Instead, they are likely to generate competing narratives about whether the program is working.

Define your use case

Establishing success criteria before a pilot begins does two things: it sets the human performance baseline the AI will be measured against. It also forces a conversation about what's actually being measured. A well-defined use case is, arguably, the most critical part of your evaluation framework.

When drafting your use case definition, be as detailed as possible and include a defined user, task, output, and a measurable quality threshold.

“The AI can analyze documents” is a capability, not a use case. → “The AI reviews lease abstractions for the acquisitions team and flags non-standard clauses for counsel” is a use case.

The more precisely a use case is scoped before an AI pilot begins, the easier it is to set the right success criteria and identify the data and integration requirements the production deployment will depend on.

How StubHub defined AI success before shipping

StubHub, an Anthropic customer, conducted rigorous A/B testing with multiple AI models before committing to a production deployment, measuring against specific resolution rate and satisfaction benchmarks rather than launching on general performance impressions. The result was a 30% reduction in support costs and response times that dropped from over 20 minutes to near-instant.

“The decision to choose Claude was entirely data-driven. We tested multiple model providers side by side, and Claude consistently delivered the best results for case resolution rates and customer satisfaction scores.”

— Timothy Addison, Engineering Org Chief of Staff, StubHub

Tie success criteria to ROI before the pilot begins

Success criteria and ROI are closely related, but organizations often treat them as separate questions addressed at separate times. The programs that scale well treat them as the same question, answered pre-pilot.

Start with high-frequency, high-value workflows where ROI compounds. A process that runs once a quarter and takes two hours to complete isn't the right candidate for an AI deployment with significant integration overhead. A process that runs hundreds of times daily is a much better candidate, as the economics scale differently.

Additionally, build a lightweight total cost of ownership model before the pilot begins: API costs, integration overhead, maintenance burden, and the

opportunity cost of not deploying. Then revisit it quarterly as AI economics evolve. Model costs have dropped significantly over the past two years and architectures built around last year's or even last month's pricing assumptions may be over-engineered for today's.

Finally, when defining your success criteria, cover both leading indicators and lagging indicators of success, as well as the lag between the two. For example, a leading indicator for value derived from a developer using [Claude Code](#) may be a decrease in PR cycle times. But that doesn't directly correspond to business value. The lagging indicator, whether that's increase in revenue, improved net revenue retention, reduced churn, or something else, should also be tracked and attributed to the new features or products being brought to market.

Establish clear ownership

Someone has to be accountable for holding the program to its success criteria, and that owner needs to be named before the pilot begins. Accenture's September 2026 [Tokenomics research](#) found that 42% of organizations rely on shared IT and finance accountability with no single owner responsible for AI costs and outcomes. Without clear ownership, it becomes difficult to measure success, make tradeoffs and maintain accountability as programs scale. That may be an individual or a small team with a named lead, but the accountability has to be unambiguous.

The owner will have decision rights over product choices, the authority to escalate blockers, and executive backing when organizational friction surfaces.


- **Decision rights** enable the owner to make calls without convening a meeting and incurring the delays that follow
- **Escalation authority** ensures that blockers requiring senior attention get resolved in days rather than months
- **Executive backing** signals to the rest of the organization that the program has institutional weight, which determines whether other functions engage as partners or as observers

Decision rights, escalation authority, and executive backing aren't interchangeable, and the absence of any one of them is enough to stall a program.

Document the governance process

Accountability at the top matters, but so does the governance structure that surrounds it. The stakeholders who will eventually approve, own, or operate the deployed system need to understand how decision-making and escalation works.

That means documenting the go/no-go process before the pilot begins:



CONSIDERATION 1: STRATEGIC OBJECTIVES, ROI, AND OWNERSHIP

who has decision rights, what criteria trigger escalation, how the program communicates wins and setbacks, and what the rollout sequence looks like. For organizations that don't define this upfront, these early-stage inputs can turn into late-stage blockers.

Broad cross-functional input can improve decision quality, but input and progress require different structures. Build a governance model that gives stakeholders a clear process to trust and a clear owner to hold accountable.

Each consideration in this guide ends with two sets of questions. 'Work out' questions require cross-functional input and don't have a single owner. 'Assign' questions are ownership decisions that a CIO or business leader must make before the program advances.

Before you move forward	
For the cross-functional team	For the CIO or business leader
Work out	Assign
What metrics define success for this use case, and how will they be measured?	Who owns evaluation design and scoring methodology?
What is the human performance baseline, and how was it established?	Who reviews the TCO model as AI economics change?
What accuracy or quality threshold justifies moving from pilot to production?	Who is the executive sponsor, and what is their active commitment?
What are the go/no-go criteria and who evaluates them?	Who owns the escalation path when progress stalls?